EE5907 Pattern Recognition

Lecture 2 Decision Theory and Classifiers

© Copyright National University of Singapore. All Rights Reserved.

Decision Theory and Classifiers:

- ✓ Bayesian Decision Theory.
- ✓ The Classification Problem.
- ✓ Bayesian Classifiers for Normal Distribution.
- ✓ Maximum A Posteriori (MAP).

© Copyright National University of Singapore. All Rights Reserved.

What is Decision Theory?

- A set of quantitative methods for reaching optimal decisions.
- Is concerned with the reasoning underlying an agent's choices, whether this is a mundane choice between taking the bus or getting a taxi, or a more far-reaching choice about whether to pursue a demanding political career.
- Layman term: *how we can decide what action should be taken that could end up with the best result*?

© Copyright National University of Singapore. All Rights Reserved.

Recall: Bayes' Rule

- "Bayesian": refer to inference methods that represent "degrees of certainty" using probability theory, and which leverage Bayes' rule, to update the degree of certainty given data.
  - *Layman term: Update our belief based on prior belief and data so that we infer something with a certain degree of uncertainty*.
- Bayes' rule: a formula for computing the probability distribution over possible values of an unknown quantity H given observed data $Y = y$.

$$p(H = h|Y = y) = \frac{p(H = h)p(Y = y|H = h)}{p(Y = y)}$$

- Follows the product rule of probability: $p(h|y)p(y) = p(h)p(y|h) = p(h, y)$
- $p(H = h|Y = y)$: **posterior distribution / conditional distribution**.

© Copyright National University of Singapore. All Rights Reserved.

Recall: Bayes' Rule: in Detail

$$p(H = h|Y = y) = \frac{p(H = h)p(Y = y|H = h)}{p(Y = y)}$$

- $p(H)$: what we know about possible values of H before we see any data – the **prior distribution** (layman terms: prior belief).
- $p(Y|H = h)$: the distribution over possible outcomes Y we expect to see if $H = h$ – the **likelihood** (layman terms: the *observed* data).
- Multiplying the prior distribution and the likelihood C unnormalized joint distribution $p(H = h, Y = y)$. Convert into a normalized distribution by dividing by $p(Y = y)$, which is denoted as the **marginal likelihood**.
- $p(H = h|Y = y)$: represents our new **belief state** (updated belief) about the possible values of H – the **posterior distribution**.

$$\text{posterior} \propto \text{prior} \times \text{likelihood}$$

© Copyright National University of Singapore. All Rights Reserved.

(Bayesian) Decision Theory: in Detail

- In decision theory, we assume the decision maker, or **agent**, has a set of possible **actions**, $\mathcal{A}$, to choose from.
- Example: doctor treat someone who may have cancer, actions $a \in \mathcal{A}$:
  - Do nothing;
  - Give patient expensive drugs with bad side effect but can save their life;
- We can compute (implicitly) the cost of treating a person (test $= 0$: cancer negative; age $= 0$: young age; action $= 1$: apply the drugs).

test	age	pr(cancer)	cost-noop	cost-drugs	action
0	0	0.01	0.84	8.00	0
0	1	0.01	0.14	8.00	0
1	0	0.80	47.73	8.00	1
1	1	0.80	7.95	8.00	0

© Copyright National University of Singapore. All Rights Reserved.

(Bayesian) Decision Theory: in Detail

test	age	pr(cancer)	cost-noop	cost-drugs	action
0	0	0.01	0.84	8.00	0
0	1	0.01	0.14	8.00	0
1	0	0.80	47.73	8.00	1
1	1	0.80	7.95	8.00	0

- Every action $a \in \mathcal{A}$ has cost and benefits, depending on the underlying state of nature $h \in \mathcal{H}$. How can we explicitly model such cost / benefit?
- Define: loss function $\ell(h, a)$ – loss incurred if action a is taken at state h.
  - The next question: how to leverage this loss?
  - Equivalently: how can we actually know quantitatively what is the best action if we compute the loss function?

© Copyright National University of Singapore. All Rights Reserved.

(Bayesian) Decision Theory: in Detail

- Define: loss function $\ell(h, a)$ – action a taken at state h.

- Compute a posterior expected loss (risk) for each possible a: (where the data is defined as x)

$$\rho(a|x) \triangleq \mathbb{E}_{p(h|x)}[\ell(h, a)] = \sum_{h \in \mathcal{H}} \ell(h, a) p(h|x)$$

- We thus define the **optimal policy** $\pi^*(x)$, which is also called the **Bayes estimator** or **Bayes decision rule** $\delta^*(x)$:

$$\pi^*(x) = \arg\min_{a \in \mathcal{A}} \rho(a|x) = \arg\min_{a \in \mathcal{A}} \mathbb{E}_{p(h|x)}[\ell(h, a)]$$

© Copyright National University of Singapore. All Rights Reserved.

Decision Theory: in Detail

- The **optimal policy** $\pi^*(x)$, or **Bayes decision rule** $\delta^*(x)$:

$$\pi^*(x) = \arg\min_{a \in \mathcal{A}} \rho(a|x) = \arg\min_{a \in \mathcal{A}} \mathbb{E}_{p(h|x)}[\ell(h, a)]$$

- If we define a **utility function** $U(h, a)$ as the desirability of each action at each state, then: $U(h, a) = -\ell(h, a)$
- The optimal policy can be re-written as:

$$\pi^*(x) = \arg\max_{a \in \mathcal{A}} \mathbb{E}_h[U(h, a)]$$

- This is defined as the **maximum expected utility principle**.
- Next, we use Bayesian decision theory to decide the optimal class label to predict given an observed input $x \in \mathcal{X}$.

© Copyright National University of Singapore. All Rights Reserved.

Classification with Zero-One Loss

- Set the following assumptions:
  - The states of nature correspond to class labels: $\mathcal{H} = \mathcal{Y} = \{1, \dots, C\}$;
  - The actions correspond to class labels: $\mathcal{A} = \mathcal{Y}$.
- Under the above assumptions, a common loss function (and a very simple one) is the **zero-one** loss $\ell(h, a) = \ell_{01}(y^*, \hat{y})$:

	$\hat{y} = 0$	$\hat{y} = 1$
$y^* = 0$	0	1
$y^* = 1$	1	0

- Or mathematically $\ell_{01}(y^*, \hat{y}) = \mathbb{I}(y^* = \hat{y})$

© Copyright National University of Singapore. All Rights Reserved.

Classification with Zero-One Loss

- The posterior expected loss $\rho(a|x) \triangleq \mathbb{E}_{p(h|x)}[\ell(h,a)]$:

$$\rho(a|x) = \rho(\hat{y}|x) = p(\hat{y} \neq y^*|x) = 1 - p(\hat{y} = y^*|x)$$

- The action that minimizes the expected loss is to choose the most probable action:

$$\pi(x) = \arg\max_{\hat{y} \in \mathcal{Y}} p(\hat{y}|x)$$

- This corresponds to the **mode** of the posterior distribution, also known as the **maximum a posteriori** or **MAP** estimate.

© Copyright National University of Singapore. All Rights Reserved.

Cost-Sensitive Classification

		Estimate 0	Estimate 1	Row sum
Truth	0	TN	FP	N
	1	FN	TP	P
Col. sum		$\hat{N}$	$\hat{P}$	

- Problem of the zero-one loss?
  - Cost of false positive (ℓ_{01}) = cost of false negative (ℓ_{10})
  - For cost-sensitive classification, the loss function would be:

$$\begin{pmatrix} \ell_{00} & \ell_{01} \\ \ell_{10} & \ell_{11} \end{pmatrix}$$

  - Set $p_0 = p(y^* = 0|x)$ and $p_1 = 1 - p_0$
  - We would choose label $\hat{y} = 0$ iff $\ell_{00}p_0 + \ell_{10}p_1 < \ell_{11}p_1 + \ell_{01}p_0$.
  - Usually, $\ell_{00} = \ell_{11} = 0$. Thus, the above simplifies to: $p_1 < \frac{\ell_{01}}{\ell_{01}+\ell_{10}}$
  - Example (Try out): if a false negative costs twice as much as false positive, then we use a decision threshold of $1/3$ before declaring a positive.

© Copyright National University of Singapore. All Rights Reserved.

Classification with Reject Option

- We may be able to say "**I don't know**" instead of returning an answer that we don't really trust – choosing the **reject** option.
  - Important in domains such as medicine where we may be risk adverse.

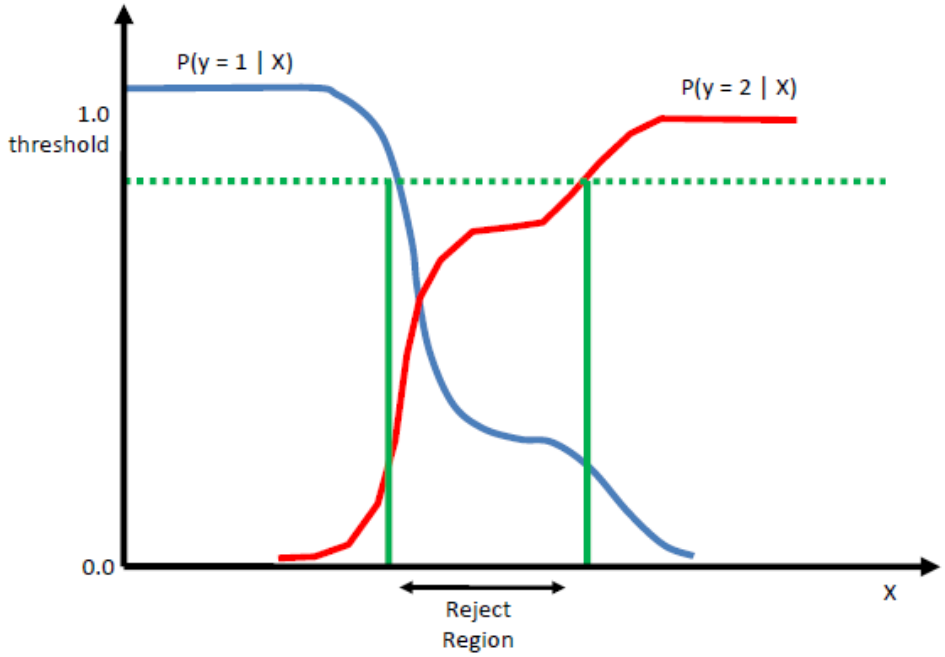

https://www.youtube.com/watch?v=P18EdAKuC1U

© Copyright National University of Singapore. All Rights Reserved.

Classification with Reject Option

- Set the following assumptions:
  - The states of nature correspond to class labels: $\mathcal{H} = \mathcal{Y} = \{1, \dots, C\}$;
  - The actions correspond to class labels: $\mathcal{A} = \mathcal{Y} \cup \{0\}$.
  - The loss function is defined as:

$$l(y^*, a) = \begin{cases} 0, & \text{if } y^* = a \text{ and } a \in \{1, \dots, C\} \\ \lambda_r, & \text{if } a = 0 \\ \lambda_e, & \text{otherwise} \end{cases}$$

  - Here λ_r: cost of the reject action, and λ_e: cost of a classification error. Set $\hat{\lambda} = 1 - \lambda_r / \lambda_e$, we thus compute the optimal policy as:

$$a^* = \begin{cases} \hat{y}, & \text{if } p^* > \hat{\lambda} \\ \text{reject}, & \text{otherwise} \end{cases}$$

  - Here $\hat{y} = \arg\max_y p(y|x)$ and $p^* = p(y^*|x) = \max_y p(y|x) = p(\hat{y}|x)$.

© Copyright National University of Singapore. All Rights Reserved.

Class Confusion Matrices

- We can choose the optimal label in a binary classification problem by **thresholding** the probability – hard threshold!
  - Instead of picking a **single threshold**, consider using a set of different thresholds, and comparing the resulting performance.
- For a threshold τ, the decision rule: $\hat{y}_\tau(x) = \mathbb{I}(p(y = 1|x) > 1 - \tau)$.
  - $FP_\tau = \sum_{n=1}^{N} \mathbb{I}(\hat{y}_n(x_n) = 1, y_n = 0)$, similar for FN_τ, TP_τ, and TN_τ.

		Predicted condition	
	Total population = P + N	Positive (PP)	Negative (PN)
Actual condition	Positive (P)	True positive (TP), hit	False negative (FN), type II error, miss, underestimation
	Negative (N)	False positive (FP), type I error, false alarm, overestimation	True negative (TN), correct rejection

© Copyright National University of Singapore. All Rights Reserved.

- The simple form of evaluation: accuracy

$$\text{Accuracy} = \frac{(TP + TN)}{(TP + FN + FP + TN)}$$

- Is this reliable?
  - If have imbalanced testing data (class 0 = 90, class 1 = 10), if the model predict **everything** as class 0, Accuracy = 90%

Class Confusion Matrices

- True positive rate (TPR), or sensitivity, recall, hit rate:

$$TPR_\tau = p(\hat{y} = 1 | y = 1, \tau) = \frac{TP_\tau}{TP_\tau + FN_\tau}$$

- False positive rate (FPR), or false alarm rate, type I error rate:

$$FPR = p(\hat{y} = 1 | y = 1, \tau) = \frac{FP_\tau}{FP_\tau + TN_\tau}$$

- Plot the TPR vs FPR as an implicit function of τ: Receiver Operating Characteristic (ROC) curve – Plot TPR vs FPR with varying τ.

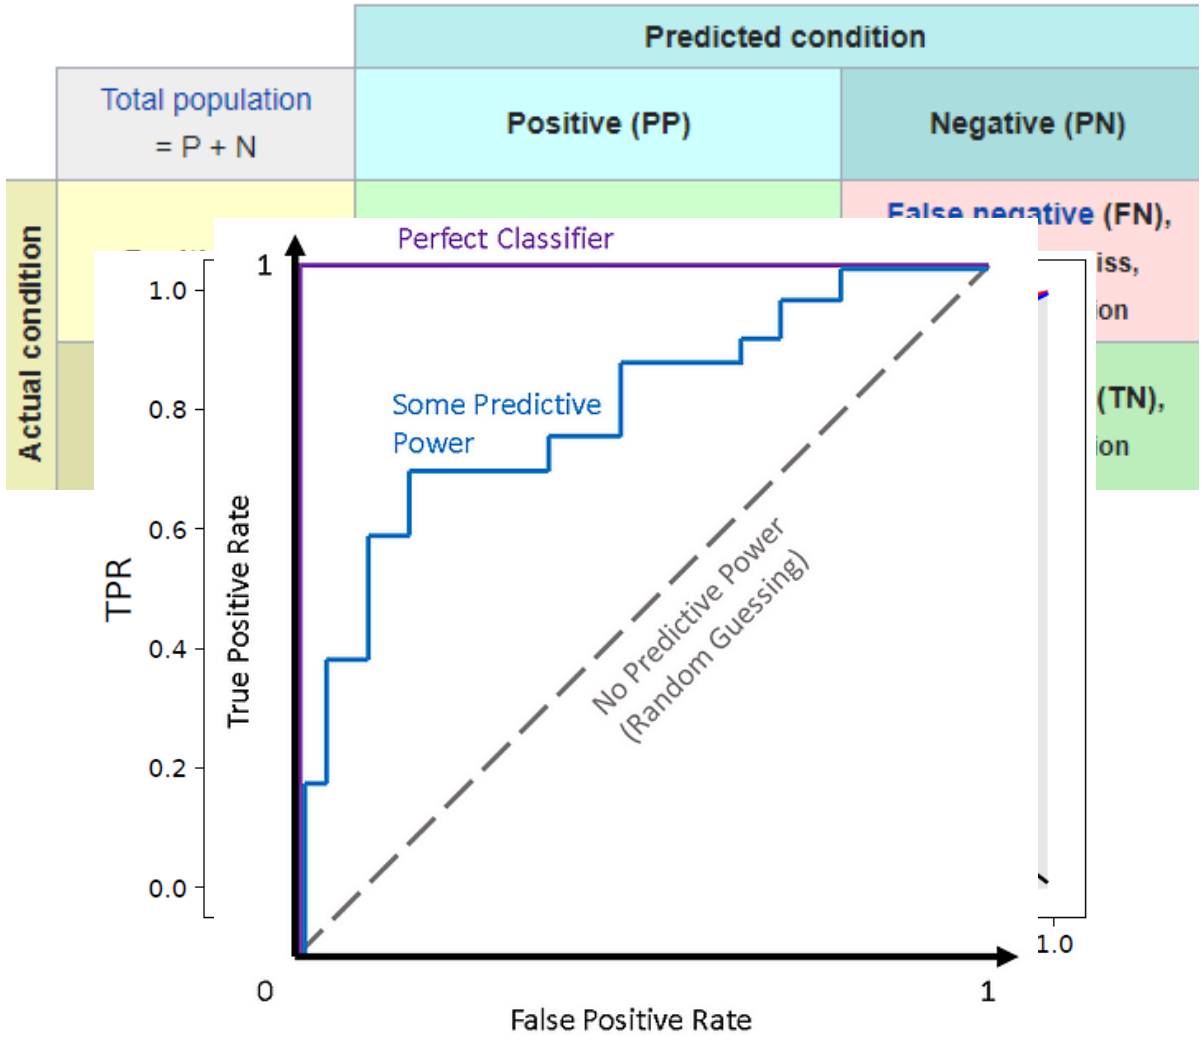

© Copyright National University of Singapore. All Rights Reserved.

End of Week 2

© Copyright National University of Singapore. All Rights Reserved.

EE5907 Pattern Recognition

CA1 Introduction

© Copyright National University of Singapore. All Rights Reserved.

Continuous Assessment #1

- Go through Assignment 1 with the instruction file.
- Deadline (strict): 3 Oct 2025.
- Late submission carry penalty marks.
- Please make sure the submitted codes are runnable, and the outputs are available when submitting.
- Individual assignment, marks will be available at later date.

© Copyright National University of Singapore. All Rights Reserved.

Recap: Class Confusion Matrices

- We can choose the optimal label in a binary classification problem by **thresholding** the probability – hard threshold!
  - Instead of picking a **single threshold**, consider using a set of different thresholds, and comparing the resulting performance.
- For a threshold τ, the decision rule: $\hat{y}_\tau(x) = \mathbb{I}(p(y = 1|x) > 1 - \tau)$.
  - $FP_\tau = \sum_{n=1}^{N} \mathbb{I}(\hat{y}_n(x_n) = 1, y_n = 0)$, similar for FN_τ, TP_τ, and TN_τ.

		Predicted condition	
	Total population = P + N	Positive (PP)	Negative (PN)
Actual condition	Positive (P)	True positive (TP), hit	False negative (FN), type II error, miss, underestimation
	Negative (N)	False positive (FP), type I error, false alarm, overestimation	True negative (TN), correct rejection

© Copyright National University of Singapore. All Rights Reserved.

- The simple form of evaluation: accuracy

$$\text{Accuracy} = \frac{(TP + TN)}{(TP + FN + FP + TN)}$$

- Is this reliable?
  - If have imbalanced testing data (class 0 = 90, class 1 = 10), if the model predict **everything** as class 0, Accuracy = 90%

Recap: Class Confusion Matrices

- True positive rate (TPR), or sensitivity, recall, hit rate:

$$TPR_\tau = p(\hat{y} = 1 | y = 1, \tau) = \frac{TP_\tau}{TP_\tau + FN_\tau}$$

- False positive rate (FPR), or false alarm rate, type I error rate:

$$FPR = p(\hat{y} = 1 | y = 1, \tau) = \frac{FP_\tau}{FP_\tau + TN_\tau}$$

- Plot the TPR vs FPR as an implicit function of τ: Receiver Operating Characteristic (ROC) curve – Plot TPR vs FPR with varying τ.

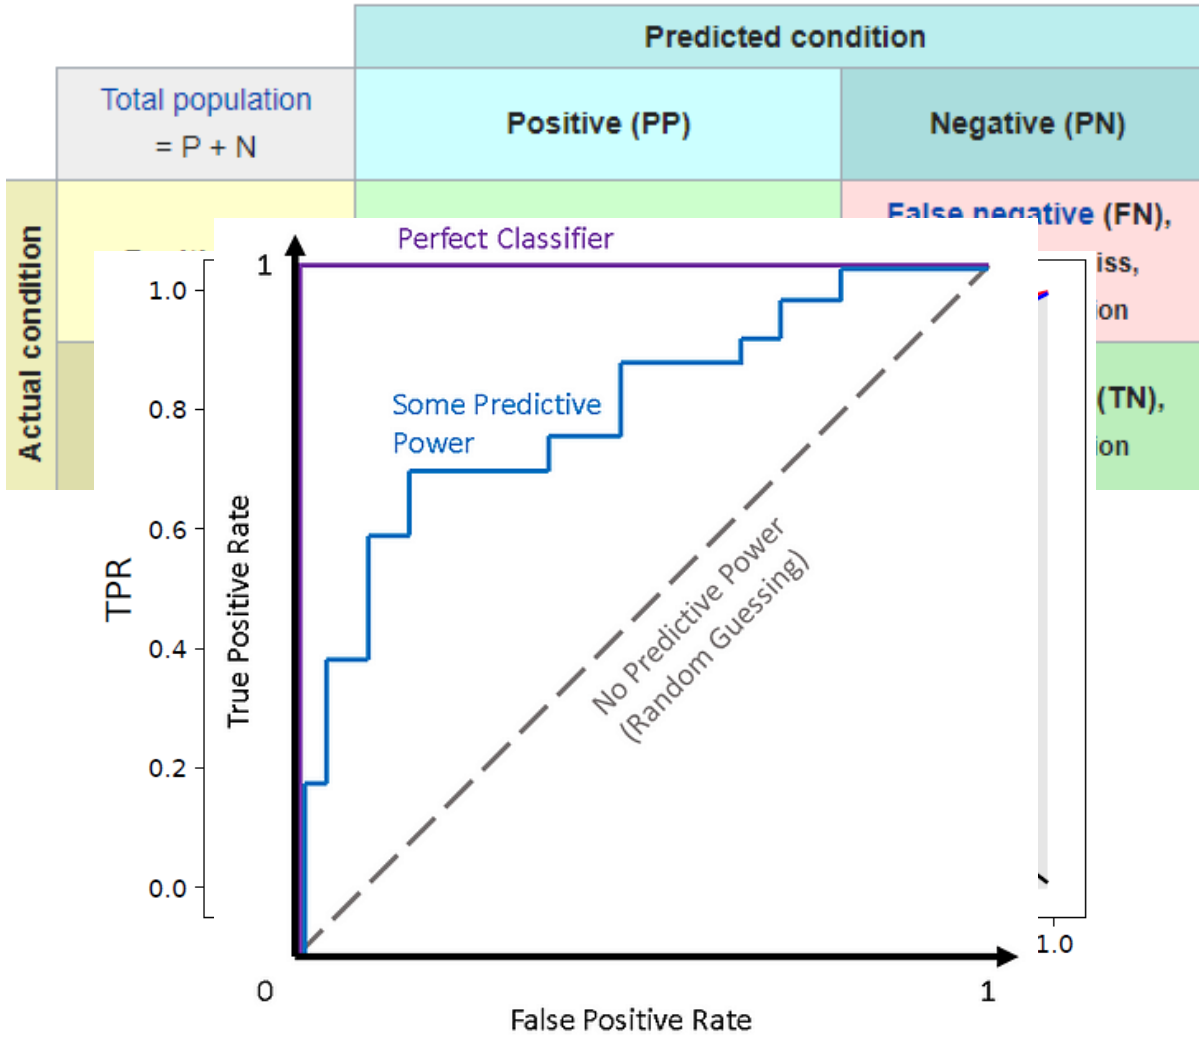

© Copyright National University of Singapore. All Rights Reserved.

Plotting the ROC Curve: Example

- ROC curve – Plot the TPR vs FPR as an implicit function of τ, i.e., plot TPR vs FPR with varying τ.
- How do we compare different ROC curves?

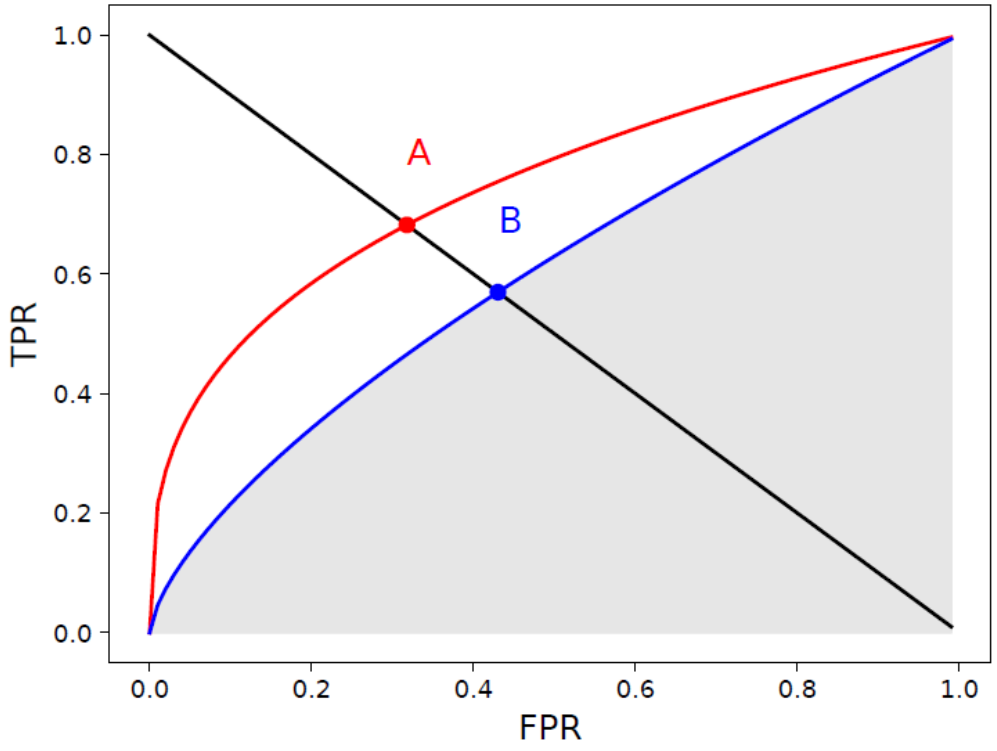

© Copyright National University of Singapore. All Rights Reserved.

The ROC Curve

- Quality of a ROC curve is often summarized using the Area Under the Curve (AUC). Higher AUC scores better; max is 1.
- Another statistic is the Equal Error Rate (EER), or the cross-over rate
  - Defined as the value which satisfies $FPR = FNR = 1 - TPR$.
  - Compute the EER by drawing a line from the top left to the bottom right and seeing where it intersects the ROC curve.
  - Lower EER scores are better; the minimum is obviously 0 (top left corner).

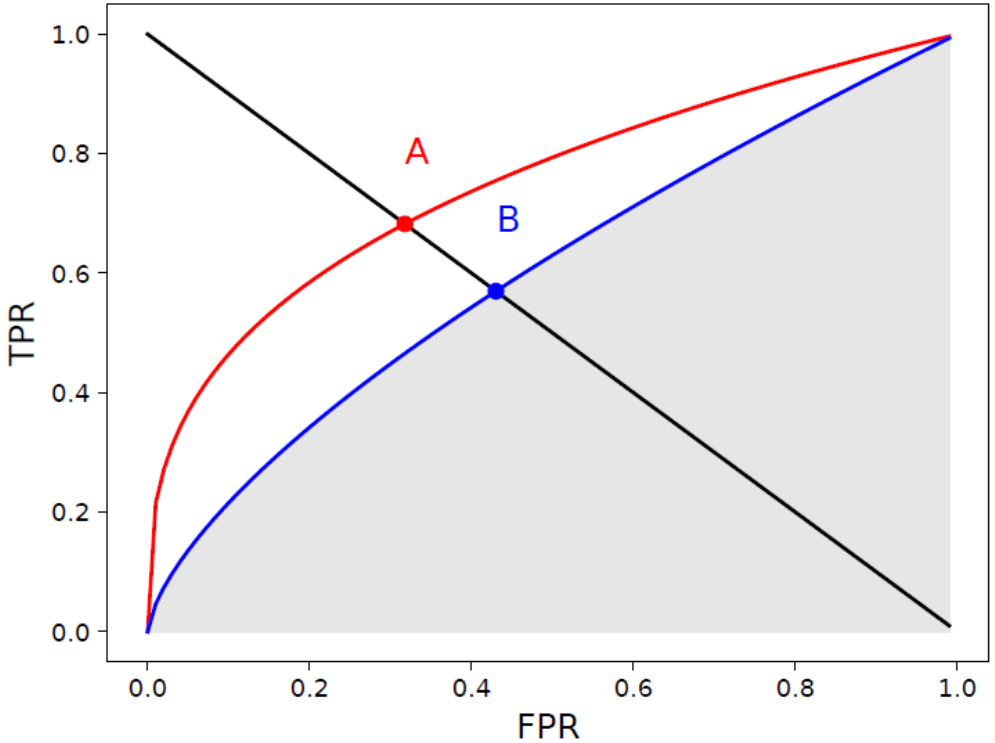

© Copyright National University of Singapore. All Rights Reserved.

Class Imbalance and the ROC Curve

- In some problems, there is severe **class imbalance**.
  - E.g., in information retrieval, the set of negatives (irrelevant items) is usually much larger than the set of positives (relevant items).
  - Discuss: would the ROC curve be affected by class imbalance?
  - Discuss: how would class imbalance affect the effectiveness of ROC?

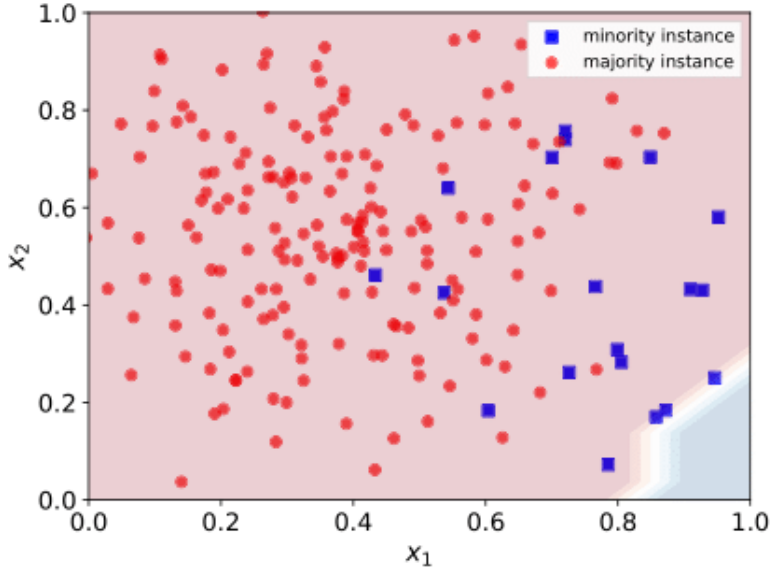

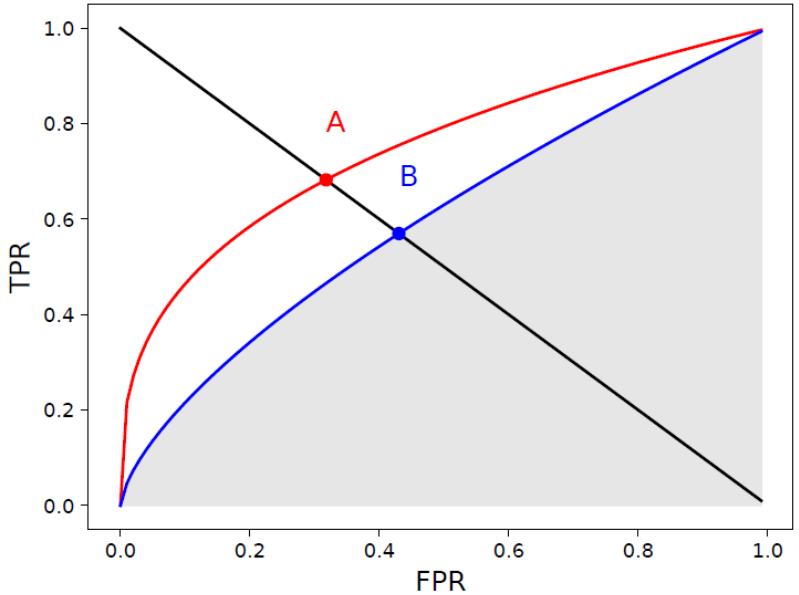

© Copyright National University of Singapore. All Rights Reserved.

The Precision-Recall Curve

- Problems of the ROC curve:
  - Require well-defined “true positives” and “true negatives”;
  - However, in some problems, “negative” may not be well-defined.
  - E.g., **detecting objects** in images: if the detector works by classifying patches, then the # of patches examined — the # of true negatives — is a param of the algorithm, instead of the problem definition.
  - Can we evaluate the classification computed just from positives?

- Precision and the precision-recall curve:
  - Precision: replace the FPR with a quantity that is computed just from positives.

© Copyright National University of Singapore. All Rights Reserved.

The Precision-Recall (PR) Curve

- Precision: fraction of our detections are actually positive:
- $\mathcal{P}(\tau) \triangleq p(y = 1 | \hat{y} = 1, \tau) = TP_\tau / (TP_\tau + FP_\tau)$
- Recall (TPR): fraction of the positives we actually detected:
- $\mathcal{R}(\tau) \triangleq p(\hat{y} = 1 | y = 1, \tau) = TP_\tau / (TP_\tau + FN_\tau)$

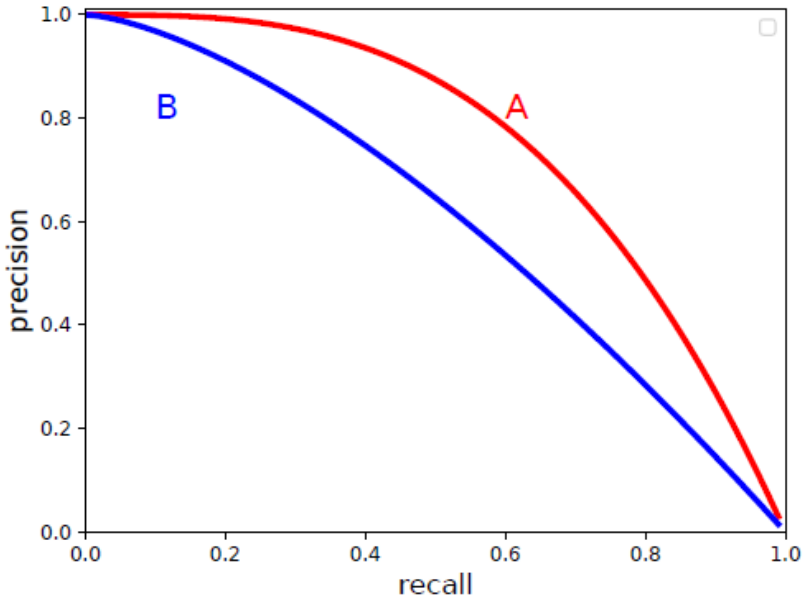

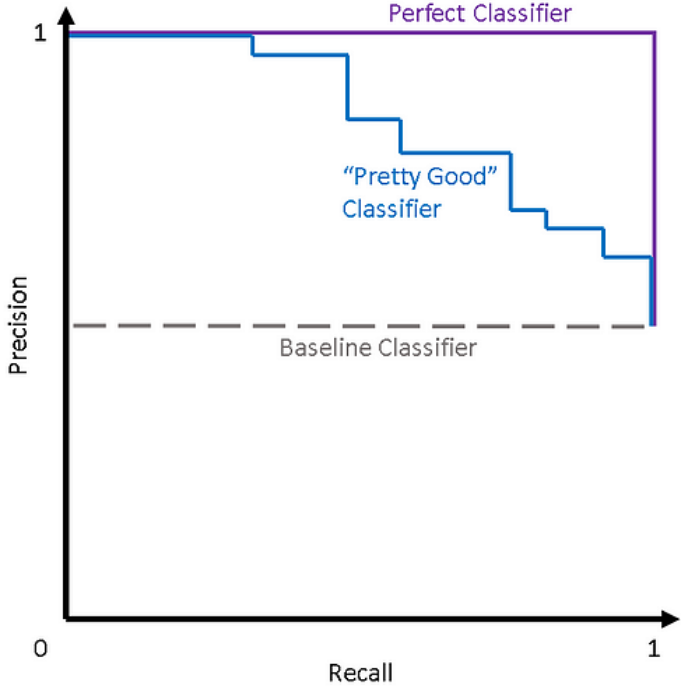

© Copyright National University of Singapore. All Rights Reserved.

An Example of PR Curve in Object Detection

- Run object detector on all test images (with NMS).
- For each **category**, compute Average Precision (AP) = area under Precision vs Recall Curve.
  - For each detection (highest score to lowest score):
  - If it matches some GT box with IoU > 0.5, mark it as positive and eliminate the GT.
  - Otherwise mark it as negative.
  - Plot a point on PR Curve.
  - Average Precision (AP) = area under PR curve
- Mean Average Precision (mAP) = average of AP for each category.
- For "COCO mAP": Compute mAP@thresh for each IoU threshold (0.5, 0.55, 0.6, ..., 0.95) and take average.

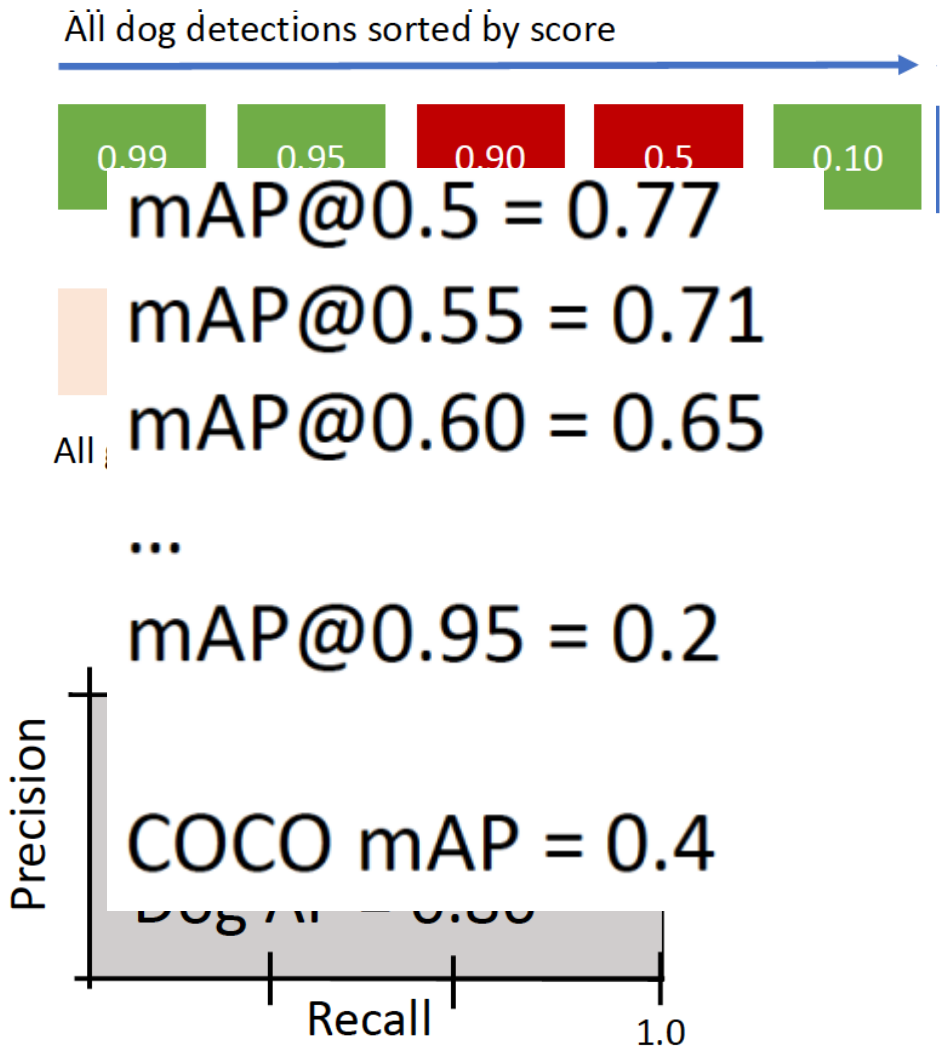

© Copyright National University of Singapore. All Rights Reserved.

Plotting the PR Curve: Example

- PR curve – Plot the precision vs recall as an implicit function of τ, i.e., plot precision vs recall with varying τ.
- How do we compare different PR curves?

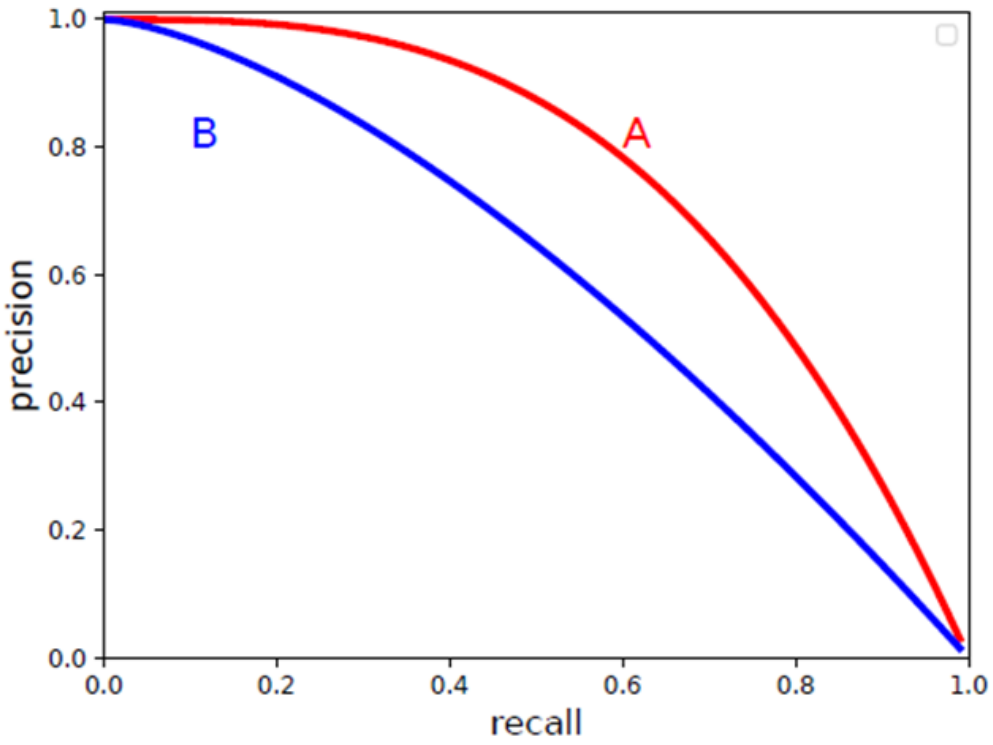

© Copyright National University of Singapore. All Rights Reserved.

The PR Curve

- The statistics to summarize the PR curve:
  - Precision at K score: the precision for a fixed recall level;
  - AUC: the area under the PR curve;
- **Interpolated precision**:
  - Possible precision **does not drop monotonically** with recall. E.g., a classifier has 90% precision at 10% recall, and 96% precision at 20% recall.
  - Measure the **maximum** precision we can achieve with **at least a recall of 10%** (which would be 96%).
- **Average precision**: avg of interpolated precisions
  - Equal to the area under the interpolated PR curve.
  - The **mean average precision** (mAP): the mean of the AP over a set of different PR curves.

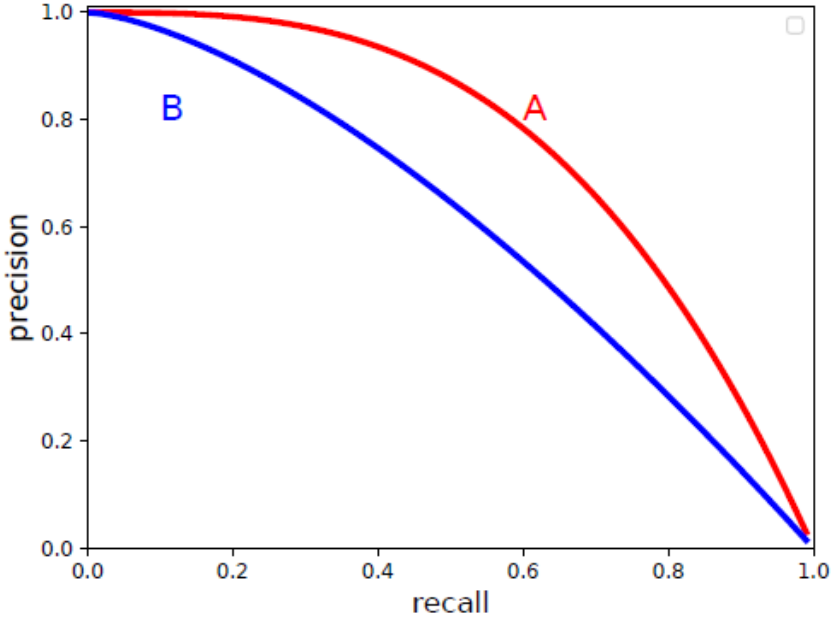

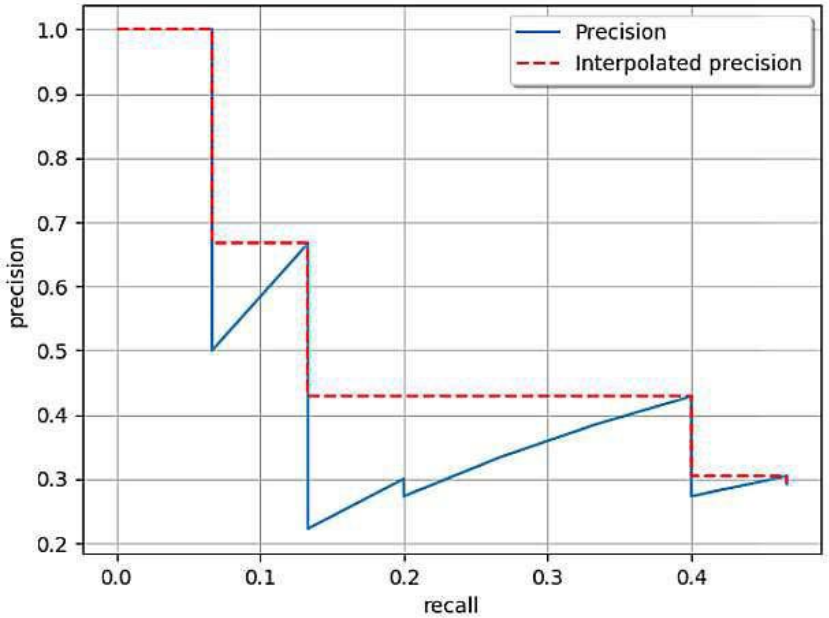

© Copyright National University of Singapore. All Rights Reserved.

The PR Curve and F-scores

- For a given point on the PR curve, can combine the precision value $\mathcal{P}$ and the recall value $\mathcal{R}$ into a single statistics:

$$\frac{1}{F_\beta} = \frac{1}{1+\beta^2}\frac{1}{\mathcal{P}} + \frac{\beta^2}{1+\beta^2}\frac{1}{\mathcal{R}}$$

- If we set $\beta = 1$, then we get the harmonic mean of the precision and recall, known as the F_1 score.

$$\frac{1}{F_1} = \frac{1}{2}\left(\frac{1}{\mathcal{P}} + \frac{1}{R}\right)$$

- Questions:
  - Why do we use the "harmonic mean" instead of the "arithmetic mean" $(\mathcal{P}+\mathcal{R})/2$?
  - Would the PR curve be affected by class imbalance?

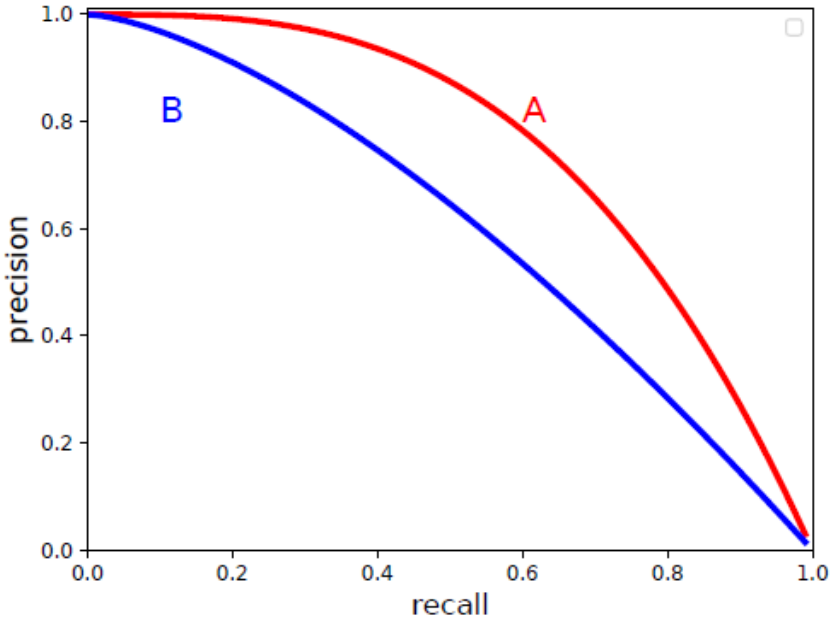

© Copyright National University of Singapore. All Rights Reserved.

Probabilistic Prediction Problem

- In previous slides: assumed the set of possible actions was to pick a single class label (or possibly the "reject" action).
- More often: the set of possible **actions** is to pick a **probability distribution** over some value of interest:
  - Perform probabilistic prediction, rather than predicting a specific value.
  - For Classification: states of nature → class labels: $\mathcal{H} = \mathcal{Y} = \{1, \ldots, C\}$; actions → class labels: $\mathcal{A} = \mathcal{Y}$.
  - For Probabilistic prediction: states of nature → $h = p(Y|x)$; action → $a = q(Y|x)$.
  - What we want: to pick q to minimize $\mathbb{E}[\ell(h, a)] = \mathbb{E}[\ell(p, q)]$ for a given x.
  - The key: how to define the loss function $\ell(p, q)$?

© Copyright National University of Singapore. All Rights Reserved.

Decision Theory and Classifiers:

- ✓ KL-Divergence
- ✓ Entropy and Mutual Information

© Copyright National University of Singapore. All Rights Reserved.

KL Divergence (KL-Div)

- Loss function is measured between two “distributions”.
- Kullback Leibler divergence (**KL divergence**) is a measure of how one probability distribution is different from a second:

$$\mathbb{D}_{KL}(p \parallel q) \triangleq \sum_{y^* \in \mathcal{Y}} p(y^*) \log \frac{p(y^*)}{q(\hat{y})}$$

- Here p is the true distribution and q is the predicted distribution.
- KL divergence can be further expanded as:
- $\mathbb{D}_{KL}(p \parallel q) = \sum p(y^*) \log p(y^*) - \sum p(y^*) \log q(\hat{y}) = -\mathbb{H}(p) + \mathbb{H}_{ce}(p, q)$. $\mathbb{H}(p)$: entropy of p; $\mathbb{H}_{ce}$: cross entropy between p and q.

© Copyright National University of Singapore. All Rights Reserved.

Entropy

- The key component of KL Divergence: entropy.
  - The **entropy** of a **probability distribution** can be interpreted as a measure of uncertainty, or lack of predictability, associated with a random variable drawn from a given distribution.
  - Entropy: define the **information content** of a data source.
  - Suppose we observe a sequence of symbols $X_n \sim p$ generated from distribution p. What does it mean by p having high / low entropy?
  - If p has high entropy, hard to predict the value of each observation X_n. Hence dataset $\mathcal{D} = (X_1, \ldots, X_n)$ has high information content.
  - If p is a degenerate distribution with 0 entropy, then every X_n will be the same, so $\mathcal{D}$ does not contain much information.

© Copyright National University of Singapore. All Rights Reserved.

Entropy for Discrete Random Variables

- More specifically, formulate the entropy of a discrete RV X with distribution p over K states (classes) as:

$$\mathbb{H}(X) \triangleq -\sum_{k}^{K} p(X=k)\log_2 p(X=k) = -\mathbb{E}_X[\log p(X)]$$

  - Note that $\mathbb{H}(X)$ is equivalent to the notation of $\mathbb{H}(p)$.
  - Question: when will the distribution obtain the Max / Min Entropy?

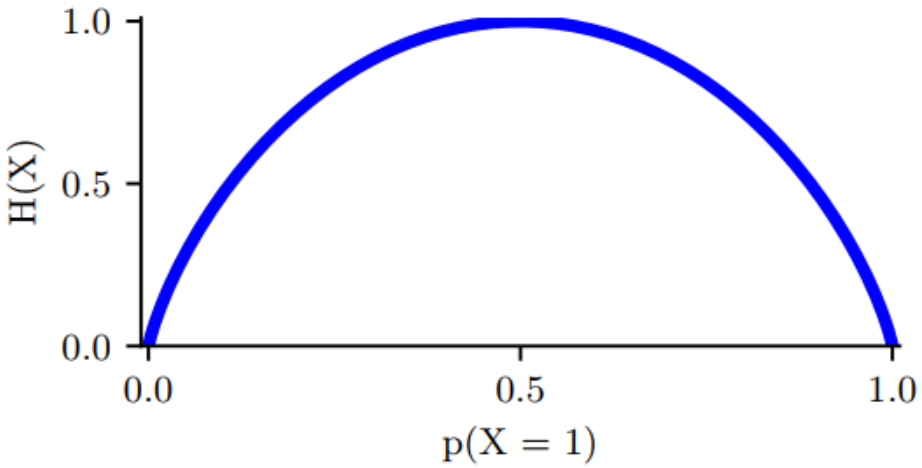

© Copyright National University of Singapore. All Rights Reserved.

Entropy for Discrete Random Variables

$$\mathbb{H}(X) \triangleq -\sum p(X=k)\log_2 p(X=k) = -\mathbb{E}_X[\log p(X)]$$

- Discrete distribution with Maximum entropy is uniform distribution.
- For a K-ary RV, the entropy is max-ed if $p(x=k) = 1/K$; $H(X) = \log_2 K$.
- Distribution with min entropy is any **delta-function** that puts all its mass on one state

- Special case: binary RVs $X \in \{0,1\}$, $p(X = 1) = \theta$ and $p(X = 0) = 1 - \theta$:
  - $\mathbb{H}(X) = -[p(X=1)\log_2 p(X=1) + p(X=0)\log_2 p(X=0)]$
  - $\mathbb{H}(X) = -[\theta \log_2 \theta + (1-\theta)\log_2(1-\theta)]$

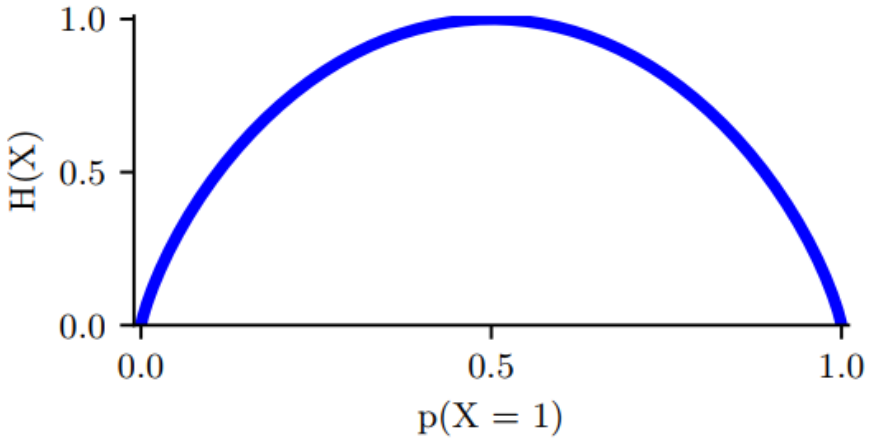

© Copyright National University of Singapore. All Rights Reserved.

Mutual Information

- KL divergence: measure how **dissimilar** two distributions were.
- Then can we measure how **dependent** two RVs are?
  - Turn the question of measuring the **dependence** of two RVs into a question about the **similarity** of their distributions.
- The notion of Mutual Information (MI) btw two RVs, defined as:

$$\mathbb{I}(X;Y) \triangleq \mathbb{D}_{KL}\big(p(x,y) \parallel p(x)p(y)\big) = \sum_{y}\sum_{x} p(x,y)\log\frac{p(x,y)}{p(x)p(y)}$$

  - Note: the notation is defined as $\mathbb{I}(X;Y)$ as X and/or Y can represent a set of variable – $\mathbb{I}(X;Y,Z)$ is the MI between X and (Y,Z).
  - MI is always non-negative: $\mathbb{I}(X;Y) \geq 0$. Reach the bound of 0 iff $p(x,y) = p(x)p(y)$.

© Copyright National University of Singapore. All Rights Reserved.

Mutual Information: Interpretation

$$\mathbb{I}(X;Y) \triangleq \mathbb{D}_{KL}\big(p(x,y) \parallel p(x)p(y)\big) = \sum_{y}\sum_{x} p(x,y)\log\frac{p(x,y)}{p(x)p(y)}$$

- MI is a KL-Div between the joint and factored marginal distributions:
- MI measures the **information gain** if we update from a model that treats the two variables as independent $p(x)p(y)$ to one that models their true joint density $p(x,y)$.
- To gain further, re-express MI in terms of joint and conditional entropies:
  $\mathbb{I}(X;Y) = \mathbb{H}(X) - \mathbb{H}(X|Y) = \mathbb{H}(Y) - \mathbb{H}(Y|X), \mathbb{H}(Y|X) = \mathbb{H}(X,Y) - H(X)$
- MI between X and Y is thus expressed as the reduction in uncertainty about X after observing Y, or, by symmetry, the reduction in uncertainty about Y after observing X.

© Copyright National University of Singapore. All Rights Reserved.

Mutual Information: Interpretation

$$\mathbb{I}(X;Y) \triangleq \mathbb{D}_{KL}\big(p(x,y) \parallel p(x)p(y)\big) = \sum_{y}\sum_{x} p(x,y)\log\frac{p(x,y)}{p(x)p(y)}$$

$$\mathbb{I}(X;Y) = \mathbb{H}(X) - \mathbb{H}(X|Y) = \mathbb{H}(Y) - \mathbb{H}(Y|X), \mathbb{H}(Y|X) = \mathbb{H}(X,Y) - H(X)$$

- An alternative proof that conditioning, on average, reduces entropy!

$$\mathbb{I}(X;Y) = \mathbb{H}(X) - \mathbb{H}(X|Y) \geq 0 \rightarrow \mathbb{H}(X|Y) \leq \mathbb{H}(X)$$

- One can also obtain the following equivalent results of MI:

$$\mathbb{I}(X;Y) = \mathbb{H}(X,Y) - \mathbb{H}(X|Y) - \mathbb{H}(Y|X)$$

$$\mathbb{I}(X;Y) = \mathbb{H}(X) + \mathbb{H}(Y) - \mathbb{H}(X,Y)$$

© Copyright National University of Singapore. All Rights Reserved.

Mutual Information: An Example

- Suppose two RVs X and Y are related to the input integer $n \in \{1, \dots, 8\}$.
  - Define $X(n) = 1$ if n is even and $Y(n) = 1$ if n is a prime number. Thus:

n	1	2	3	4	5	6	7	8
X	0	1	0	1	0	1	0	1
Y	0	1	1	0	1	0	1	0

$p(X,Y)$	$Y=0$	$Y=1$
$X=0$	$\frac{1}{8}$	$\frac{3}{8}$
$X=1$	$\frac{3}{8}$	$\frac{1}{8}$

	Y=0	Y=1
X=0	$\frac{1}{4}$	$\frac{3}{4}$
X=1	$\frac{3}{4}$	$\frac{1}{4}$

  - $\mathbb{H}(X,Y) = -\left[\frac{1}{8}\log\frac{1}{8} + \frac{3}{8}\log\frac{3}{8} + \frac{3}{8}\log\frac{3}{8} + \frac{1}{8}\log\frac{1}{8}\right] = 1.81\ bits$
  - $\mathbb{H}(X) = \mathbb{H}(Y) = 1$
  - $\mathbb{H}(Y|X) \triangleq -\sum p(X,Y)\log p(Y|X) = -\left[\frac{1}{8}\log\frac{1}{4} + \frac{3}{8}\log\frac{3}{4} + \frac{3}{8}\log\frac{3}{4} + \frac{1}{8}\log\frac{1}{4}\right] = 0.81$
  - Verify that $\mathbb{H}(Y|X) = \mathbb{H}(X,Y) - H(X) = 0.81$
  - $\mathbb{I}(X;Y) = \mathbb{H}(X) - \mathbb{H}(X|Y) = 1 - 0.81 = 0.19\ bits$

© Copyright National University of Singapore. All Rights Reserved.

Entropy and Mutual Information

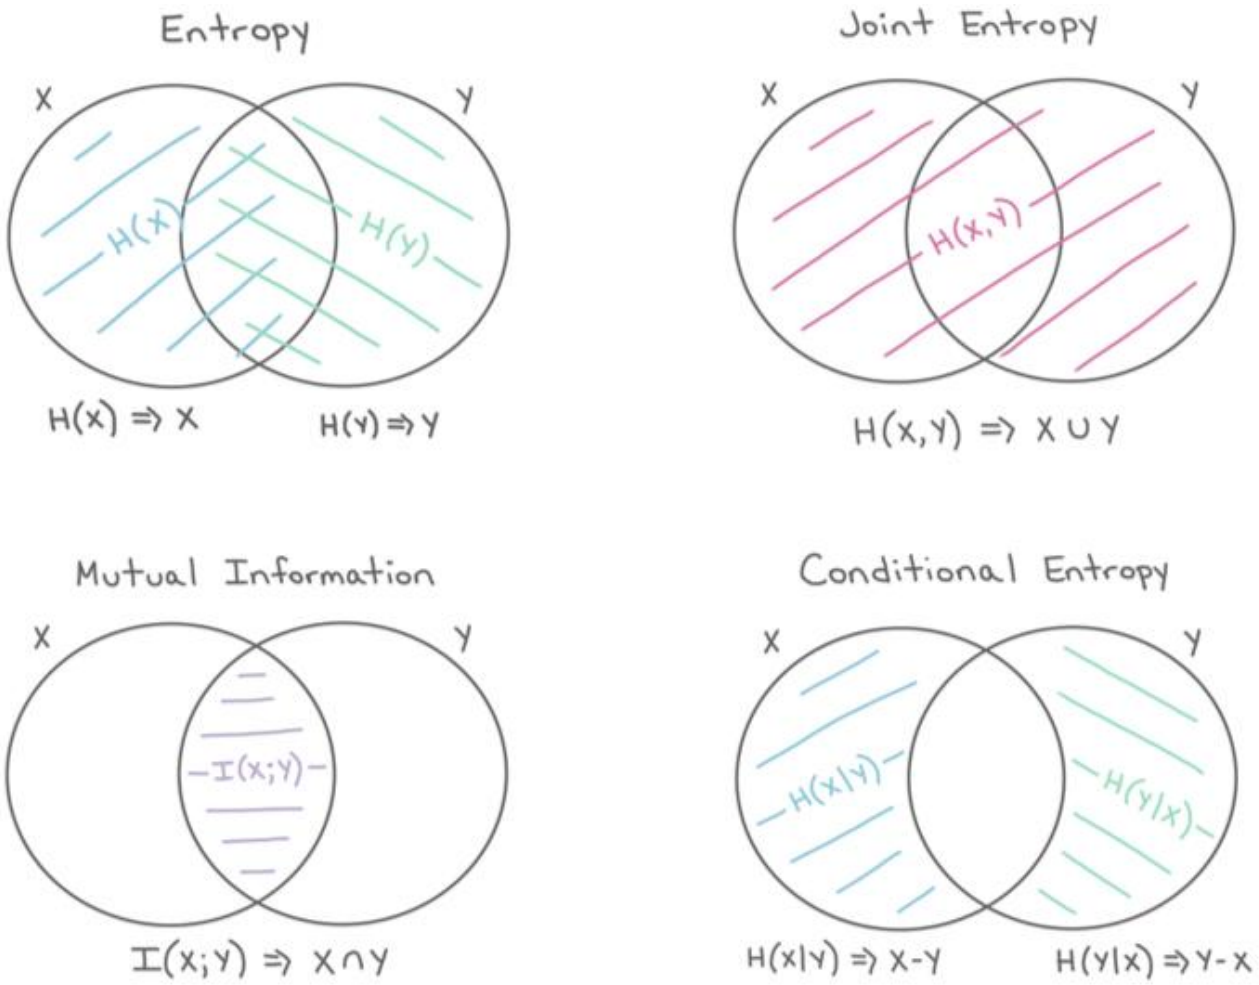

© Copyright National University of Singapore. All Rights Reserved.

Decision Theory and Classifiers:

- ✓ The Frequentist Decision
- ✓ Application in the Context of Supervised Learning

© Copyright National University of Singapore. All Rights Reserved.

Recall: Frequentist vs. Bayesian

- Frequentist interpretation: probabilities represent long run frequencies of events that can happen multiple times.
- E.g., Flipping the coin many times, we expect it to land heads about half the time.
- Bayesian interpretation: quantify **uncertainty** or ignorance about something
- E.g., The coin is equally likely to land heads/tails on next toss.

© Copyright National University of Singapore. All Rights Reserved.

Frequentist Decision Theory

- In the **Frequentist decision theory**, treat the **unknown** state of nature (denoted by $\boldsymbol{\theta}$ instead of h) as a **fixed yet unknown quantity**, and treat the data x as random.
  - Instead of conditioning over x, we **average** over it to compute the loss expect to incur if apply our estimator to **different datasets**.

- Define the frequentist **risk** of an estimator δ given an unknown state of nature $\boldsymbol{\theta}$ as the expected loss when applying δ to data x, where the expectation is over the data, sampled from $p(x|\boldsymbol{\theta})$:

$$R(\boldsymbol{\theta}, \delta) \triangleq \mathbb{E}_{p(x|\boldsymbol{\theta})}\left[\ell\left(\boldsymbol{\theta}, \delta(x)\right)\right]$$

© Copyright National University of Singapore. All Rights Reserved.

The Bayes Risk

$$R(\boldsymbol{\theta}, \delta) \triangleq \mathbb{E}_{p(x|\boldsymbol{\theta})}\left[\ell\left(\theta, \delta(x)\right)\right]$$

- In general, the true state of nature $\boldsymbol{\theta}$ that generates the data x is unknown, so cannot compute the risk directly. Then how?
- One solution: to assume a prior π_0 for θ then average it out – the **Bayes Risk** (integrated risk):

$$R_{\pi_0}(\delta) \triangleq \mathbb{E}_{\pi_{0(\boldsymbol{\theta})}}[R(\boldsymbol{\theta}, \delta)] = \int d\boldsymbol{\theta} dx \pi_0(\boldsymbol{\theta}) p(x|\boldsymbol{\theta}) \ell\left(\boldsymbol{\theta}, \delta(x)\right)$$

- Decision rule that minimize the Bayes risk: **Bayes estimator**.

© Copyright National University of Singapore. All Rights Reserved.

The Bayes Estimator

- Recall: in the Bayesian decision theory, we define the **optimal policy**, also the **Bayes estimator** or **Bayes decision rule** as:

$$\pi^*(x) = \arg\min_{a \in \mathcal{A}} \rho(a|x) = \arg\min_{a \in \mathcal{A}} \mathbb{E}_{p(h|x)}[\ell(h, a)]$$

- In the Frequentist decision theory, the optimal estimator is:

$$\delta(x) = \arg\min_{a} \int d\theta \pi_0(\theta) p(x|\theta) \ell(\theta, a) \equiv \arg\min_{a} \int d\theta p(\theta|x) \ell(\theta, a)$$

- The two forms are actually equivalent! What does this mean?
  - Bayesian approach provides a good way of achieving the frequentist goal – picking the optimal action case-by-case is optimal on average!

© Copyright National University of Singapore. All Rights Reserved.

The Maximum Risk

- Use of a prior might seem undesirable in certain cases. We can therefore define the **maximum risk** as:

$$R_{\max}(\delta) \triangleq \sup_{\boldsymbol{\theta}} R(\boldsymbol{\theta}, \delta)$$

- Decision rule that min the max risk: minimax estimator δ_{MM}.
- One example on the right:
  - δ_1 has lower worst-case risk than δ_2.
  - δ_1 is the minimax estimator (even if δ_2 has lower risk for most values of $\boldsymbol{\theta}$)/
  - The minimax estimators are overly conservative!

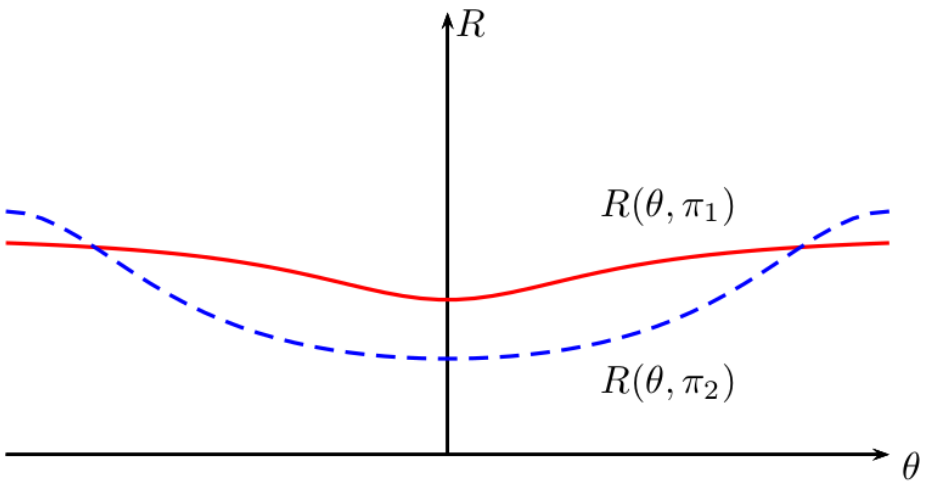

© Copyright National University of Singapore. All Rights Reserved.

ERM and Frequentist Decision Theory

- We consider how to apply the frequentist decision theory in the context of **supervised learning**.
- The risk: $R(\boldsymbol{\theta}, \delta) \triangleq \mathbb{E}_{p(x|\boldsymbol{\theta})}\left[\ell\left(\boldsymbol{\theta}, \delta(x)\right)\right]$.
  - Unknown "state of nature" can correspond to the unknown **parameters** of some model $\boldsymbol{\theta}^*$.
  - Meanwhile, the discussion w.r.t. data within a dataset $x \in \mathcal{D}$, the risk can be rewritten:

$$R(\delta, \boldsymbol{\theta}^*) = \mathbb{E}_{p(\mathcal{D}|\theta^*)}\left[\ell\left(\boldsymbol{\theta}^*, \delta(\mathcal{D})\right)\right]$$

© Copyright National University of Singapore. All Rights Reserved.

ERM and Frequentist Decision Theory

- In supervised learning,
  - Different unknown state of nature (output y) for each different input x,
  - Estimator δ is a prediction function $\hat{y} = f(x)$,
  - The state of nature (as a whole) is the true distribution $p^*(x, y^*)$,
  - Thus, the risk of an estimator is (in terms of f and p^*) is:
  - $R(f, p^*) = R(f) \triangleq \mathbb{E}_{p^*(x)p^*(y|x)}[\ell(y^*, f(x))]$
  - Defined as the **population risk**, as the expectations are taken w.r.t. the true joint distribution $p^*(x, y^*)$.
- Can we compute the **population risk**? If not how?

© Copyright National University of Singapore. All Rights Reserved.

ERM and Frequentist Decision Theory

- In most cases, quite impossible to obtain the unknown true distribution p^*.
  - Instead, can approximate it with the empirical distribution of N samples:
  - $p_{\mathcal{D}}(x, y|\mathcal{D}) \triangleq \frac{1}{|\mathcal{D}|}\sum_{(x_n,y_n)\in\mathcal{D}} \delta(x - x_n)\delta(y - y_n)$
  - Here the empirical distribution is exactly the training distribution $p_{\mathcal{D}}(x, y|\mathcal{D}) = p_{tr}(x, y)$. We thus obtain the empirical risk (which is a random variable) as:

$$R(f, \mathcal{D}) \triangleq \mathbb{E}_{p_{\mathcal{D}}(x, y|\mathcal{D})}[\ell(y, f(x)] = \frac{1}{N}\sum_{n=1}^{N} \ell\big(y_n, f(x_n)\big)$$

  - A natural way to choose the predictor is to choose the predictor with minimum risk:
  - $\hat{f}_{ERM} = \arg\min_f R(f, \mathcal{D}) = \arg\min_f \frac{1}{N}\sum_n^N \ell(y_n, f(x_n)) = \arg\min_\theta \mathcal{L}(\theta)$
  - $\mathcal{L}(\boldsymbol{\theta}) = \sum_{n=1}^N \ell(y_n, \boldsymbol{\theta}; x_n)$ (as defined in Lecture 1)
- The process of choosing predictor $\hat{f}_{ERM}$ is **empirical risk minimization (ERM)**.

© Copyright National University of Singapore. All Rights Reserved.

Performance of Functions Fitted with ERM

- What is the theoretical performance of functions that are fitted with the ERM principle?
  - Let $f^{**} = \arg\min_{f} R(f)$ be the function that achieves the min possible population risk, where we optimize over ALL possible functions.
  - !! **ALL possible functions**?
  - Define $f^{*} = \arg\min_{f \in \mathcal{H}} R(f)$ the best function in hypothesis space $\mathcal{H}$.
  - We still cannot compute what f^{*} is – the population risk R is unknown!
  - We can only compute by further approx. – minimizing the empirical risk:

$$f_N^{*} = \arg\min_{f \in \mathcal{H}} R\left(f, \mathcal{D}\right) = \arg\min_{f \in \mathcal{H}} \mathbb{E}_{p_{tr}}\left[\ell\left(y, f(x)\right)\right]$$

© Copyright National University of Singapore. All Rights Reserved.

Performance of Functions Fitted with ERM

- The risk of our chosen predictor f_N^* compared to the best possible predictor f^{**} can be decomposed into two terms:
- $\mathbb{E}_{p^*}[R(f_N^*) - R(f^{**})] = \underbrace{R(f^*) - R(f^{**})}_{\varepsilon_{app}(\mathcal{H})} + \underbrace{\mathbb{E}_p[R(f_N^*) - R(f^*)]}_{\varepsilon_{est}(\mathcal{H}, N)}$
  - $\varepsilon_{app}(\mathcal{H})$ is the approximate error: measures how closely $\mathcal{H}$ can model the true optimal function f^{**}.
  - $\varepsilon_{est}(\mathcal{H}, N)$ is the estimation error or generalization error: measures the difference in estimated risk due to having a finite training set.
  - Approx. $\varepsilon_{est}(\mathcal{H}, N)$ by the difference btw training set and test set error: $\varepsilon_{est}(\mathcal{H}, N) \approx \mathbb{E}_{p_{tr}}\left[\ell\left(y, f_N^*(x)\right)\right] - \mathbb{E}_{p_{te}}\left[\ell\left(y, f_N^*(x)\right)\right]$ – Generalization Gap.

© Copyright National University of Singapore. All Rights Reserved.

Regularization and Regularized Risk

- To avoid the chance of overfitting, it is common to add a complexity penalty to the objective function – regularization.
  - The action that minimizes the expected loss is to choose the most probable action:

$$\pi(x) = \arg\max_{y \in \mathcal{Y}} p(y|x)$$

  - This corresponds to the **mode** of the posterior distribution, also known as the **maximum a posteriori** or **MAP** estimate.
  - If the loss function is a log loss term, while the regularizer is a negative log prior, $R_\lambda(\theta, \mathcal{D}) = -\frac{1}{N}\sum_{n=1}^{N} \log p(y_n|x_n, \theta) - \lambda \log p(\theta)$. Minimizing this is exactly equivalent to the MAP estimation.

© Copyright National University of Singapore. All Rights Reserved.

Occam's Razor

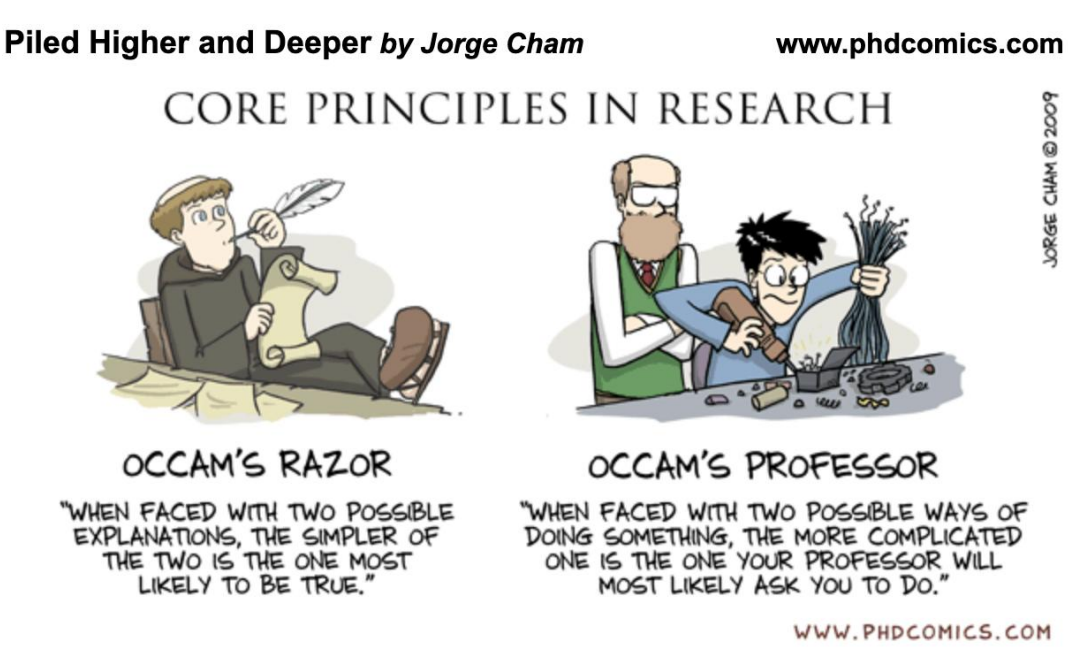

- **Occam's razor** is the problem-s
  searching for explanations cons
  set of elements.
  - Also known as the **principle of parsimony** or the **law of parsimony.**
  - Consider 2 models, a simple one, m_1, and a more complex one, m_2. Suppose both can explain the data by suitably optimizing their parameters, i.e., $p(\mathcal{D}|\widehat{\boldsymbol{\theta}}_1, m_1)$ and $p(\mathcal{D}|\widehat{\boldsymbol{\theta}}_2, m_2)$ are both large.
- Intuitively should prefer m_1, as it is **simpler** and as good as m_2.
- In regularization, it corresponds to the complexity penalty, where more complex models would result in a larger regularization term.

© Copyright National University of Singapore. All Rights Reserved.

Population Risk Est. with Validation

- A simple way to estimate the population risk for a supervised learning setup: cross-validation.
  - Partition the dataset into two, one used for training models (training set), another, the **validation** set or **holdout** set, used for assessing the risk.
  - Note that in supervised learning, the real "test data" are **inaccessible**.
- Process: fit the model on the training set, use its performance on the validation set as an approx. to the population risk:
  - Make the dependence of the $R_\lambda(f, \mathcal{D})$ on the dataset more explicit:

$$R_\lambda(\boldsymbol{\theta}, \mathcal{D}) = \frac{1}{|D|} \sum_{(x,y)\in\mathcal{D}} \ell\big(y, f(x; \boldsymbol{\theta})\big) + \lambda C(\boldsymbol{\theta})$$

  - Also define $\widehat{\boldsymbol{\theta}}_\lambda(\mathcal{D}) = \arg\min_\theta R_\lambda(\theta, \mathcal{D})$, $\mathcal{D}_{tr}$ and $\mathcal{D}_{val}$ are partitions.

© Copyright National University of Singapore. All Rights Reserved.

Cross-Validation

- For each model, we fit it to the training set to obtain $\widehat{\boldsymbol{\theta}}_\lambda(\mathcal{D}_{tr})$. Then use the **unregularized** empirical risk on the validation set as an estimate of the population risk – validation risk: $R_\lambda^{val} \triangleq R_0\left(\widehat{\boldsymbol{\theta}}_\lambda(\mathcal{D}_{tr}), \mathcal{D}_{val}\right)$.
  - Any shortcomings of this
- Consider the training d
  Impossible to make reli
- Solution?
- Cross-validation!

© Copyright National University of Singapore. All Rights Reserved.

Cross-Validation

- Split the training data into K folds;
- For each fold $k \in \{1, \ldots, K\}$, train on all but ONE fold, in a round-robin fashion (in turns, circled around all possible folds).

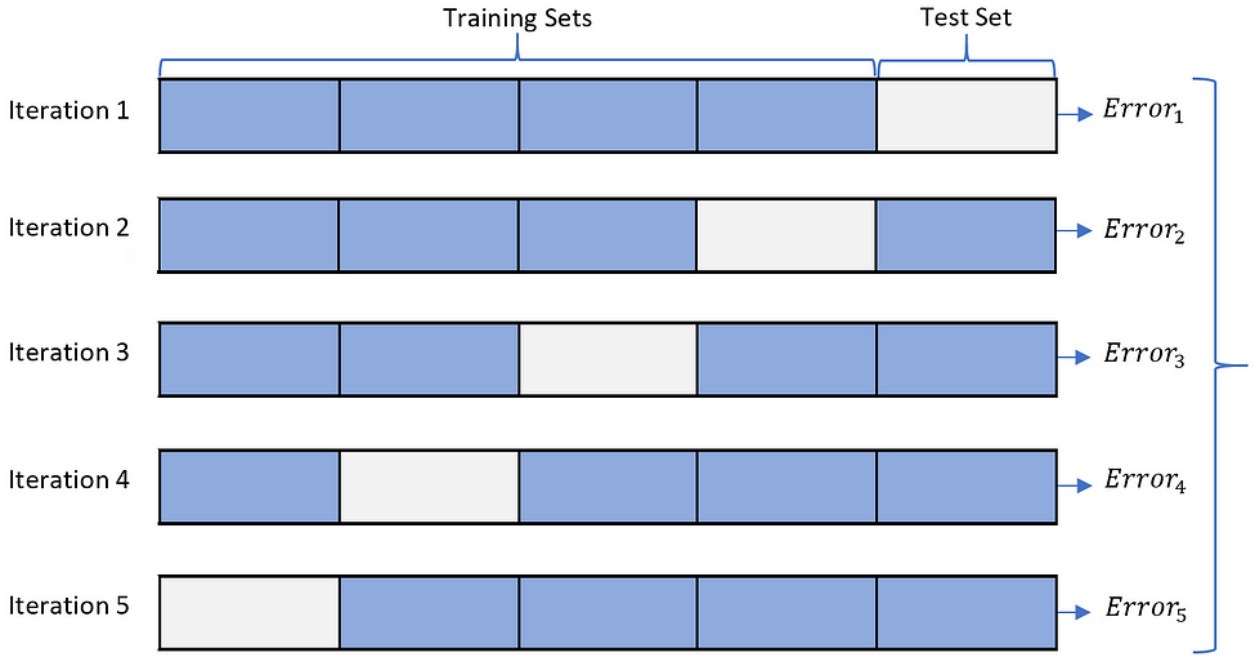

- Formally, the cross-validation risk is formulated as:
- $R_\lambda^{cv} \triangleq \frac{1}{K}\sum_{k=1}^{K} R_0\left(\widehat{\boldsymbol{\theta}}_\lambda(\mathcal{D}_{-k}), \mathcal{D}_k\right)$
- $\mathcal{D}_k$ is the data in the k-th fold, and $\mathcal{D}_{-k}$ is all the other data. If set $K = N$, the method is known as leave-one-out cross-validation.
- Alternatively, we can estimate an optimal $\hat{\lambda}$ with CV, then re-estimate the model parameters using all data: $\widehat{\boldsymbol{\theta}} = \arg\min_{\boldsymbol{\theta}} R_{\hat{\lambda}}(\boldsymbol{\theta}, \mathcal{D})$.

© Copyright National University of Singapore. All Rights Reserved.

Cross-Validation: Example

```python
from sklearn.model_selection import KFold
k = 5
kf = KFold(n_splits=k, shuffle=True, random_state=0)
```

✓ 0.0s

```python
from sklearn.model_selection import cross_val_score
scores_sklearn = cross_val_score(model, X, y, cv=kf, scoring='accuracy')
mean_accuracy_sklearn = np.mean(scores_sklearn)

print(f"Accuracy Score for each fold: {[float(score) for score in scores_sklearn]}")
print(f"Mean accuracy across {k} folds: {mean_accuracy_sklearn}")
```

✓ 0.0s

© Copyright National University of Singapore. All Rights Reserved.